

HPC Final

Vu Dinh Anh - m21.ict.001

January 2023

Figure 1 and Figure 2 show the differences between before and after applying filter with window size of 8 ($7 + 1$ including the target pixel)




Figure 1: Original image




Figure 2: Kuwahara filtered

This code was copied from previous labwork probably 8.

```

@cuda.jit
def rgb_to_hsv(src, dst):
    i = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
    j = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
    r, g, b = src[i, j, 0], src[i, j, 1], src[i, j, 2]
    max_c = max(r, g, b)
    min_c = min(r, g, b)
    d = max_c - min_c

    # h
    if d == np.float64(0):
        dst[i, j, 0] = np.float64(0)
    if max_c == r:
        dst[i, j, 0] = 60 * ((g - b)/d % 6)
    if max_c == g:
        dst[i, j, 0] = 60 * ((b - r)/d + 2)
    if max_c == b:
        dst[i, j, 0] = 60 * ((r - g)/d + 4)

    # s
    if max_c == np.float64(0):
        dst[i, j, 1] = np.float64(0)
    if max_c != np.float64(0):
        dst[i, j, 1] = d/max_c

    # v
    dst[i, j, 2] = max_c

```

An end-to-end function is used to package and show the step. It's not much but honest work.

```

def e2e_rgb_to_hsv(img_in):
    img_in = np.array(img_in, copy=True)
    img_in = np.float64(img_in)
    img_in /= 255
    img_in = np.ascontiguousarray(img_in)
    h, w, c = img_in.shape
    hsv = np.array(img_in, copy=True)

    # Configure CUDA blocks
    block_size = (32, 32)
    grid_size = dual_tuple_division(block_size, (h, w))

    cuda_img_in = cuda.to_device(img_in)
    cuda_hsv = cuda.to_device(hsv)

    rgb_to_hsv[grid_size, block_size](cuda_img_in, cuda_hsv)
    hsv = cuda_hsv.copy_to_host()
    return hsv

```

I implemented Kuwahara filter poorly. Each window's standard deviation is implemented manually. Then I find minimum standard deviation. Lastly I check min std manually. It's not a good solution.

```

@cuda.jit
def kuwahara_filter(rgb_src, rgb_dst, v_arr, window_size):
    i = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x

```

```

j = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
height, width = v_arr.shape

# window 0
sum_ = np.float64(0)
sum_power_2 = np.float64(0)
for wi in range(i - window_size, i):
    for wj in range(j - window_size, j):
        if wi >= 0 and wi < height and wj >= 0 and wj < width:
            sum_ += v_arr[wi, wj]
            sum_power_2 += v_arr[wi, wj] * v_arr[wi, wj]

mean_ = sum_ / (window_size * window_size)
std_window_0 = math.sqrt(sum_power_2 / (window_size * window_size) - mean_ * mean_)

# window 1
sum_ = np.float64(0)
sum_power_2 = np.float64(0)
for wi in range(i, i + window_size):
    for wj in range(j - window_size, j):
        if wi >= 0 and wi < height and wj >= 0 and wj < width:
            sum_ += v_arr[wi, wj]
            sum_power_2 += v_arr[wi, wj] * v_arr[wi, wj]

mean_ = sum_ / (window_size * window_size)
std_window_1 = math.sqrt(sum_power_2 / (window_size * window_size) - mean_ * mean_)

# window 2
sum_ = np.float64(0)
sum_power_2 = np.float64(0)
for wi in range(i - window_size, i):
    for wj in range(j, j + window_size):
        if wi >= 0 and wi < height and wj >= 0 and wj < width:
            sum_ += v_arr[wi, wj]
            sum_power_2 += v_arr[wi, wj] * v_arr[wi, wj]

mean_ = sum_ / (window_size * window_size)
std_window_2 = math.sqrt(sum_power_2 / (window_size * window_size) - mean_ * mean_)

# window 3
sum_ = np.float64(0)
sum_power_2 = np.float64(0)
for wi in range(i, i + window_size):
    for wj in range(j, j + window_size):
        if wi >= 0 and wi < height and wj >= 0 and wj < width:
            sum_ += v_arr[wi, wj]
            sum_power_2 += v_arr[wi, wj] * v_arr[wi, wj]

mean_ = sum_ / (window_size * window_size)
std_window_3 = math.sqrt(sum_power_2 / (window_size * window_size) - mean_ * mean_)

# find min std
min_std = min(std_window_0, std_window_1, std_window_2, std_window_3)

```

```

# assign avg rgb of window to the target pixel
avg_r = np.float64(0.0)
avg_g = np.float64(0.0)
avg_b = np.float64(0.0)
if min_std == std_window_0:
    # window 0
    for wi in range(i- window_size, i):
        for wj in range(j - window_size, j):
            if wi >= 0 and wi < height and wj >= 0 and wj < width:
                avg_r += rgb_src[wi, wj, 0]
                avg_g += rgb_src[wi, wj, 1]
                avg_b += rgb_src[wi, wj, 2]

elif min_std == std_window_1:
    # window 1
    for wi in range(i, i + window_size):
        for wj in range(j - window_size, j):
            if wi >= 0 and wi < height and wj >= 0 and wj < width:
                avg_r += rgb_src[wi, wj, 0]
                avg_g += rgb_src[wi, wj, 1]
                avg_b += rgb_src[wi, wj, 2]
elif min_std == std_window_2:
    # window 2
    for wi in range(i - window_size, i):
        for wj in range(j, j + window_size):
            if wi >= 0 and wi < height and wj >= 0 and wj < width:
                avg_r += rgb_src[wi, wj, 0]
                avg_g += rgb_src[wi, wj, 1]
                avg_b += rgb_src[wi, wj, 2]
else:
    # window 3
    for wi in range(i, i + window_size):
        for wj in range(j, j + window_size):
            if wi >= 0 and wi < height and wj >= 0 and wj < width:
                avg_r += rgb_src[wi, wj, 0]
                avg_g += rgb_src[wi, wj, 1]
                avg_b += rgb_src[wi, wj, 2]

avg_r = np.uint8(avg_r / (window_size * window_size))
avg_g = np.uint8(avg_g / (window_size * window_size))
avg_b = np.uint8(avg_b / (window_size * window_size))

rgb_dst[i, j, 0] = avg_r
rgb_dst[i, j, 1] = avg_g
rgb_dst[i, j, 2] = avg_b

```

0.1 Without shared memory

```

runtime @ block size
0.00025529870006266717 @ (2, 2)
0.00025514400001611646 @ (4, 4)
0.00026532719998613176 @ (8, 8)
0.00026966600003106576 @ (16, 16)

```

0.00024621869999919 @ (32, 32)

0.2 With shared memory

I'm too tired and exhausted. Therefore, I haven't got time to implement this. Some ways are:

- load rgb source image, or rgb destination image or the v channel of rgb-2-hsv converted image
- the value of the v channel after raised power to 2